




Chapter 12: Binary Search Trees


What is a Binary Search Tree (BST)?

A binary search tree is a binary tree with the following two properties:

1. If a node y is in the left subtree of a node x , then $y.key \leq x.key$.
2. If a node y is in the right subtree of a node x , then $y.key > x.key$.



(a)



(b)

Inorder tree walk can print out all the keys in a binary search tree in a sorted order.

```
INORDER-TREE-WALK(x)
1. if x != NIL
2.   INORDER-TREE-WALK(x.left)
3.   print x
4.   INORDER-TREE-WALK(x.right)
```

We can similarly do a **preorder** and **postorder** walk. All of them take worst-case runtime of $\Theta(n)$.

Recommended Exercises: 12.1-1, 12.1-2, 12.1-4. *Challenging:* 12.1-3 (solved in class)

Other Recommended Exercises:

1. Write a recursive procedure for counting the number of nodes in a binary search tree.
2. Write a recursive procedure for finding the height of a binary search tree.

Solution 12.1-3:

ITERATIVE-INORDER-TREE-WALK(*x*)

```
1. initialize an empty stack S
2. while true
3.     //traverse to the leftmost leaf
4.     while x != NIL
5.         PUSH(S, x)
6.         x = x.left
7.     //if stack is empty, then we are done
8.     if EMPTY-STACK(S)
9.         return
10.    //pop the top element, print it and then add nodes
11.    //in the right subtree to the stack
12.    x = POP(S)
13.    print x.data
14.    x = x.right
```

Example:




Step	Output	Stack
0		F,B,A
1	A	F,B
2	B	F,D,C
3	C	F,D
4	D	F,E
5	E	F
6	F	G
7	G	I,H
8	H	I
9	I	

Querying a Binary Search Tree


Querying a Binary Search Tree: retrieve information from the tree without modifying the tree.

Query operations of a binary search tree include Search, Minimum, Maximum, Successor, Predecessor, ...


All of the above operations take $O(h)$, where h is the height of the tree.




(a)



(b)



(c)



(d)

```
TREE-SEARCH(x, k)
1. if x == NIL or k == x.key
2.   return x
3. if k < x.key
4.   return TREE-SEARCH(x.left, k)
5. else return TREE-SEARCH(x.right, k)
```

The nodes searched during the recursion form a path from the root downward. Thus, running time is $O(h)$.

```

ITERATIVE-TREE-SEARCH(x, k)
1. while x != NIL and k != x.key
2.     if k < x.key
3.         x = x.left
4.     else x = x.right
5. return x

```

```

TREE-MINIMUM(x)
1. while x.left != NIL
2.     x = x.left
3. return x

```

```

TREE-MAXIMUM(x)
1. while x.right != NIL
2.     x = x.right
3. return x

```

The successor of a node x is the node y , where

$$y = \begin{cases} \text{Minimum}(\text{right}[x]) & \text{if } \text{right}[x] \neq \text{NIL} \\ \text{The lowest ancestor of } x \text{ whose left child is also an ancestor of } x & \text{if } \text{right}[x] = \text{NIL} \end{cases}$$

```

TREE-SUCCESSOR(x)
1. if x.right != NIL
2.     return TREE-MINIMUM(x.right)
3. else
4.     y = x.p
5.     while y != NIL and x == y.right
6.         x = y
7.         y = y.p
8.     return y

```

The worst-case runtime is $O(h)$, since we either follow a path downward (1st case) or a path upward (2nd case)

The predecessor of a node x is the node y , where

$$y = \begin{cases} \text{Maximum}(\text{left}[x]) & \text{if } \text{left}[x] \neq \text{NIL} \\ \text{The lowest ancestor of } x \text{ whose right child is also an ancestor of } x & \text{if } \text{left}[x] = \text{NIL} \end{cases}$$

Solution to Exercise 12.2-3:

```
TREE-PREDECESSOR(x)
1. if x.left != NIL
2.   return TREE-MAXIMUM(x.left)
3. else
4.   y = x.p
5.   while y != NIL and x == y.left
6.     x = y
7.     y = y.p
8.   return y
```

The worst-case runtime is $O(h)$, using same argument as for TREE-SUCCESSOR(x)

Recommended Exercises: 12.2-1, 12.2-2, 12.2-4, 12.2-5, 12.2-6.


Solve in class: 12.2-3, 12.2-6.

Insertion and Deletion

Insertion

TREE-INSERT(T, z): insert a node z into a binary search tree T , where $z.key = v$, $z.left = z.right = z.p = \text{NIL}$ initially.

TREE-INSERT always inserts a new node z as a leaf node.



```

TREE-INSERT(T, z)
// Insert node z, where z.key = v, z.left = NIL z.right = NIL
1.  x = T.root           // x keeps track of the path for insertion
2.  y = NIL              // y will track the parent of x
3.  while x != NIL
4.      y = x
5.      if z.key < x.key
6.          x = x.left
7.      else x = x.right
8.  z.p = y
9.  if y == NIL
10.     T.root = z
11. elseif z.key < y.key
12.     y.left = z
13. else y.right = z

```

Steps 3 - 7: find the position to insert the new node.

Steps 8 - 13: set the pointers to insert the new node.

Takes $O(h)$ worst-case runtime: trace downward from the root to a leaf to find the position to insert.

Give an example, using the animation in the references.

Questions:


- What kind of BST do we get if we insert n elements that are already sorted in ascending order? What is its height?
- What kind of BST do we get if we insert n elements that are already sorted in ascending order? What is its height?

Deletion

To delete a node z from a binary search tree, there are 3 cases to consider.

1. If z does not have any children, then modify the parent $z.p$: replace z with NIL as $z.p$'s child.
2. If z has only one child, then take out z by making a new link between its child and its parent.
3. If z has two children, then take out z 's successor y (y has no left child) and copy the contents of y to z .

These get refined into four cases in the code, shown visually below:




```

TRANSPLANT(T, u, v)
//replace the subtree rooted at node u with the subtree rooted at node v
1.  if u.p == NIL
2.      T.root = v
3.  elseif u == u.p.left
4.      u.p.left = v
5.  else u.p.right = v
6.  if v != NIL
7.      v.p = u.p

TREE-DELETE(T, z)
1.  if z.left == NIL
2.      TRANSPLANT(T, z, z.right)
3.  elseif z.right == NIL
4.      TRANSPLANT(T, z, z.left)
5.  else y = TREE-MINIMUM(z.right)
6.      if y != z.right
7.          TRANSPLANT(T, y, y.right)
8.          y.right = z.right
9.          y.right.p = y
10. TRANSPLANT(T, z, y)
11. y.left = z.left
12. y.left.p = y

```

Takes $O(h)$ worst-case runtime: case 1 or 2 take $\Theta(1)$, but case 3 takes $O(h)$.

Recommended Exercises: 12.3-1, 12.3-2, 12.3-3, 12.3-4, 12.3-5.

Solve in class: 12.3-3, 12.3-5

Randomly built binary search tree have an expected height of $O(\lg n)$, similar to the best-case.

How can we balance an existing binary search tree?

Recommended Exercise: We can do an inorder walk and then recursively build back a balanced tree. Develop pseudo-code for this procedure.

References

- TREE visualization and interactive explorer: <https://visualgo.net/en/bst>